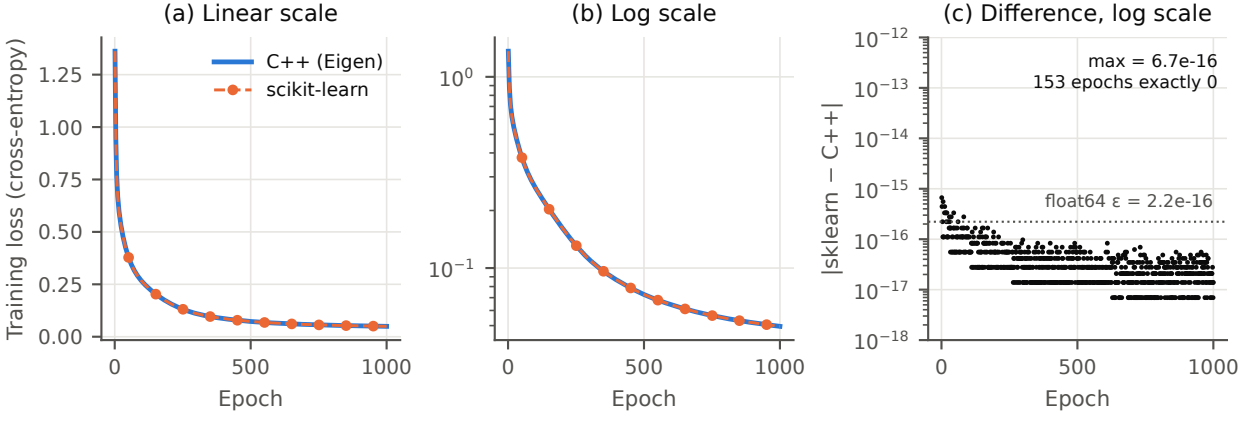
(a) Linear scale
(b) Log scale
(c) Difference, log scale
C++ (Eigen)
scikit-learn
Training loss (cross-entropy)
|sklearn − C++|
max = 6.7e-16
153 epochs exactly 0
float64 ε = 2.2e-16
Epoch
Epoch
Epoch