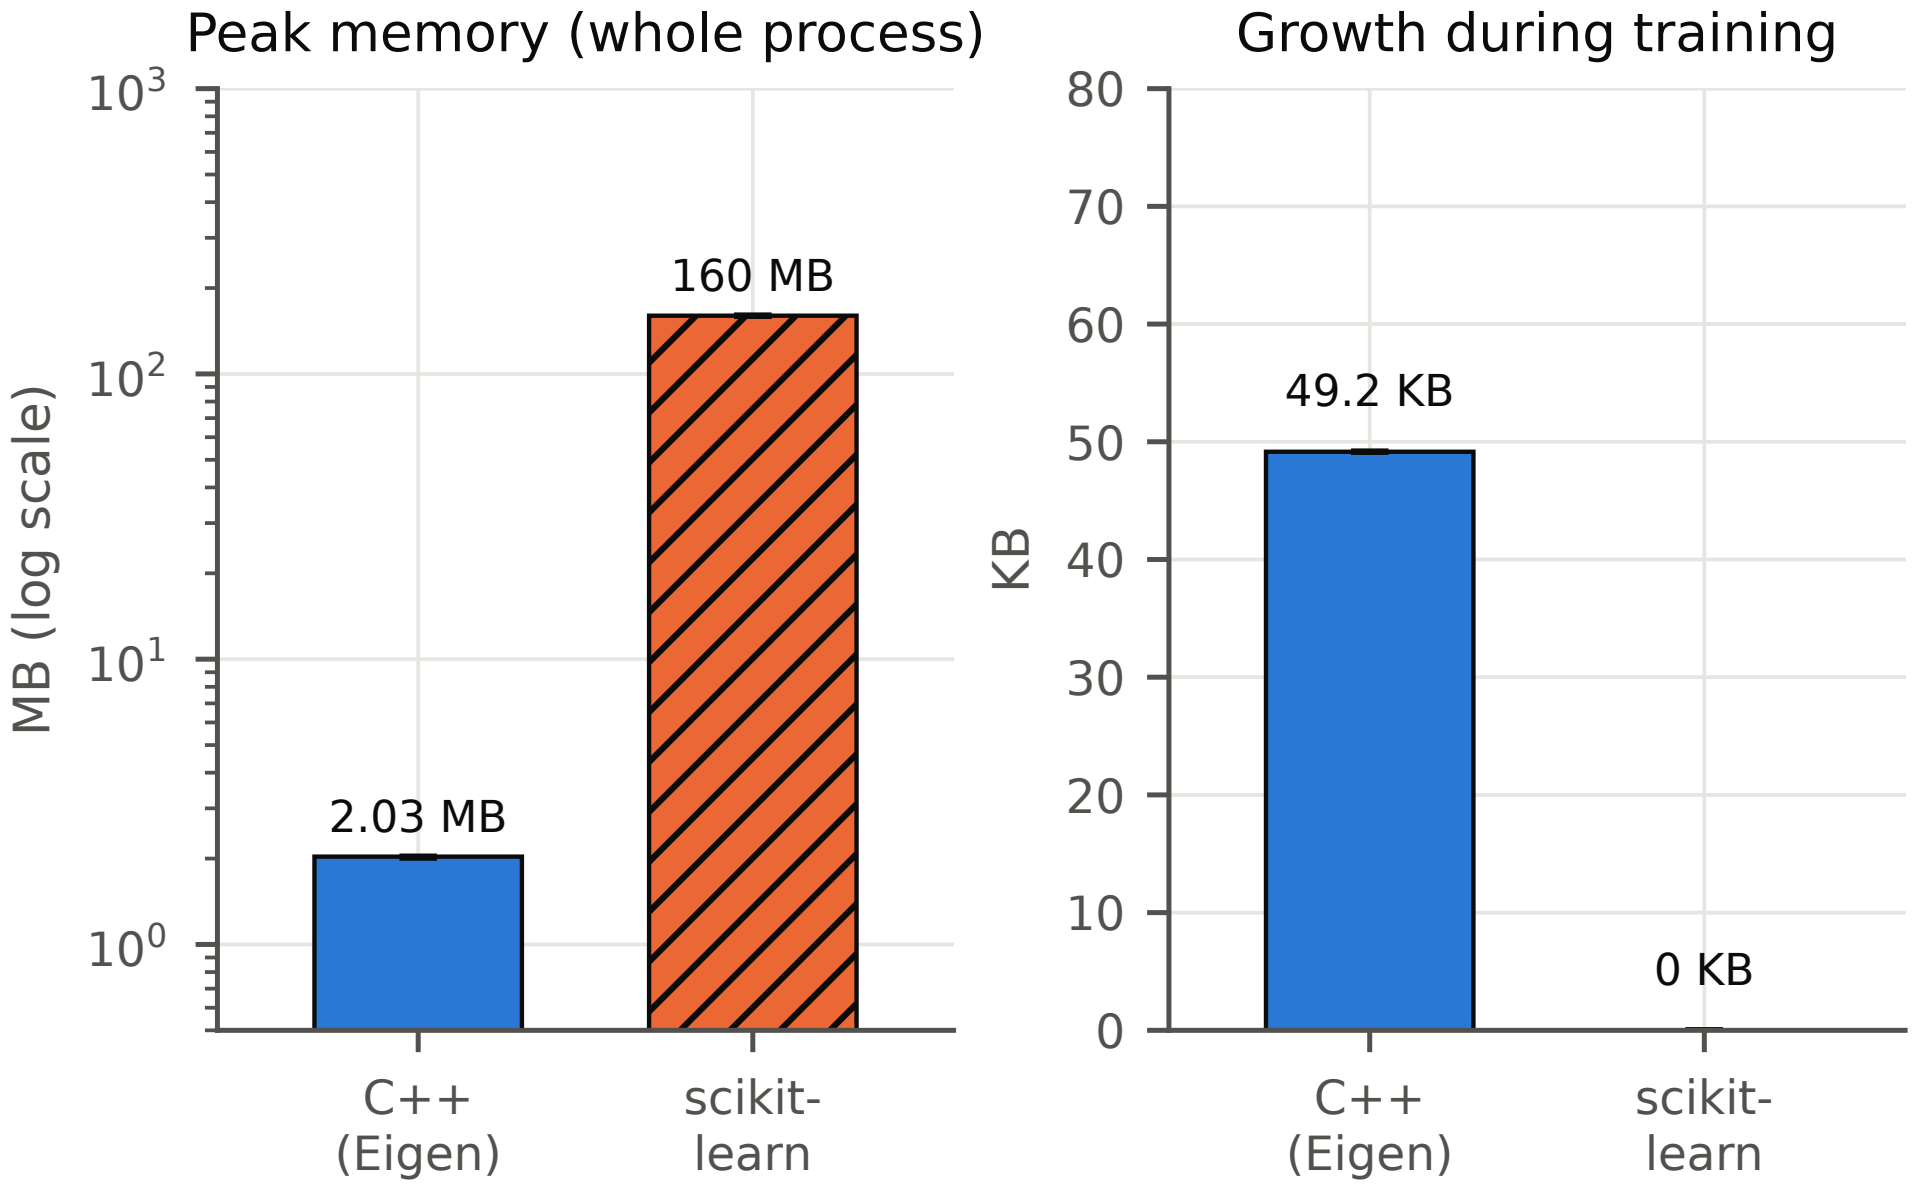

Peak = maximum resident set size of the whole process: for C++ the binary, Eigen and the data; for scikit-learn the Python interpreter, numpy, scipy, scikit-learn and pandas (scikit-learn 1.9.1 imports pandas itself, so a pandas-free peak could not be measured). Growth = increase of that peak while training (RSS is counted in 16 KB pages, so 0 means no new page). Median and min–max of 5 clean processes per side.